

Junior Cloud Security Engineer

Internship Task Report

Intern Full Name	[REDACTED]
Intern Email	[REDACTED]
Cohort	Cohort 4, 2026
Task ID	CSE-AWS-L-011-04
Task Title	Amazon S3 Exposure Detection, Policy Analysis, and Access Enumeration
Skill Domain	AWS Cloud Security / IAM / S3 Access Control / Threat Detection / STS / IaC Security
Difficulty Level	Advanced
CPE Credits	4 CPE Credits
Guiding Framework(s)	MITRE ATT&CK / NIST CSF / NIST SP 800-53 / CIS Benchmarks / ISO 27001 / OWASP / NICE Framework
Submission Date	29 Aug 2026

1.0 Task Overview

Task ID	CSE-AWS-L-011-04
Task Title	Amazon S3 Exposure Detection, Policy Analysis, and Access Enumeration
Skill Domain	AWS Cloud Security / IAM / S3 Access Control / Threat Detection / STS / IaC Security
Difficulty	Advanced
Duration	Approximately 8 hours
CPE Credits	4 CPE Credits
Cloud Platform(s)	AWS lab account; CloudGoat; Cybr; PwnedLabs
AWS Region(s) Used	AWS account region used for the S3 lab environment
Primary Tools	AWS CLI v2, CloudGoat 2.5.0, Steampipe, jq, AWS IAM/S3 Console
Target Environment	AWS lab account with S3 buckets; CloudGoat cloud_breach_s3; Cybr and PwnedLabs authorised training labs
MITRE ATT&CK IDs	T1530 (Data from Cloud Storage Object); T1087 (Account Discovery); T1078 (Valid Accounts)
Guiding Frameworks	MITRE ATT&CK; NIST CSF; NIST SP 800-53; CIS AWS Foundations Benchmark; ISO/IEC 27001
Regulatory Alignment	ISO/IEC 27001 access control and logging principles; NIST SP 800-53 AC-3, AC-6, AU-2; CIS S3 public-access guidance

1.1 Introduction

This task examined how Amazon S3 exposure can occur through public ACLs, permissive bucket policies, disabled Block Public Access controls, and excessive IAM permissions. The work combined policy analysis, AWS CLI enumeration, CloudGoat deployment, and authorised hosted labs to understand how exposed cloud storage can be discovered and abused, and how least privilege and public-access controls reduce that risk.

1.2 Objectives & Learning Outcomes

Understand how Amazon S3 exposure occurs through misconfigurations such as public ACLs, overly permissive bucket policies, and disabled Block Public Access settings — Met by reviewing S3 bucket-policy structure and analysing Effect, Principal, Action, and Resource, including the risk of wildcard principals.

Deploy a misconfigured S3 environment using CloudGoat — Met by preparing the Amazon Linux 2023 EC2 instance, installing/configuring CloudGoat and dependencies, configuring the cloudgoat AWS CLI profile, and validating the S3 environment.

Enumerate accessible S3 buckets and analyse ACLs, policies, principals, and public-access controls — Met by using AWS CLI, jq, and Steampipe to identify bucket configurations, permissions, and AccessDenied conditions.

Extend practical S3 enumeration skills through authorised hosted labs — Met by completing **Cybr – Introduction to S3 Enumeration** and **PwnedLabs – Identify the AWS Account ID from a Public S3 Bucket**.

Document findings, risks, and mitigations in a professional cloud-security report — Met by consolidating evidence, analysing the security impact of S3 exposure, mapping relevant MITRE ATT&CK concepts, and recommending S3 Block Public Access and least-privilege IAM controls.

1.3 Scope & Legal Authorisation

All activities were performed only within authorised lab environments. No production systems, external AWS accounts, third-party resources, or real customer data were intentionally accessed. AI assistance was used for explanation, troubleshooting guidance, risk-analysis wording, and report drafting; technical commands and evidence were executed and captured in the authorised lab environment.

Authorised environments used:

- AWS lab account — S3 test environment and IAM user [REDACTED].
- CloudGoat — cloud_breach_s3 scenario on the EC2 lab instance.
- PwnedLabs — Identify the AWS Account ID from a Public S3 Bucket.
- Cybr — Introduction to S3 Enumeration.

No external AWS accounts, production environments, third-party cloud resources, or real customer data were accessed. AI was used only for learning support, troubleshooting explanations, analysis, and documentation; it did not execute AWS commands or create lab screenshots.

1.4 MITRE ATT&CK Cloud Threat Mapping

MITRE ATT&CK Technique ID	Technique Name	Relevance to This Task
T1530	Data from Cloud Storage Object	Directly relevant to the risk of retrieving data from an exposed or publicly accessible S3 bucket.
T1087	Account Discovery	Relevant to enumeration and discovery activities used to understand accessible cloud resources and account context.
T1078	Valid Accounts	The [REDACTED] IAM credentials and configured AWS profile demonstrate how valid credentials determine which S3 enumeration actions are authorised.

2.0 Environment & Tool Setup

2.1 Cloud Architecture Overview

Cloud Resource / Component	Type / Service	Role in Task	Region / Config Notes
EC2 CloudGoat host	Amazon EC2 / Amazon Linux 2023	Hosted CloudGoat, AWS CLI, Python virtual environment, jq and Steampipe	Private lab instance; CloudGoat working directory under ec2-user
[REDACTED]	AWS IAM user	Authenticated identity used for S3 enumeration and policy/ACL testing	Configured in AWS CLI profile cloudgoat
secure-s3-lab-sk	Amazon S3	Primary S3 test bucket used for policy and ACL analysis	Bucket policy allowed the [REDACTED] principal broad S3 access during the lab
Other lab S3 buckets	Amazon S3	Enumeration targets visible to the lab account	Included CloudTrail/config/lab buckets returned by aws s3 ls
CloudGoat cloud_breach_s3	CloudGoat scenario	Purpose-built vulnerable S3 training scenario	Authorised sandbox scenario

2.2 AWS CLI & Tool Configuration

Tool / Service	Version	Installation Method	Purpose in This Task
AWS CLI	2.33.15	Preinstalled on Amazon Linux 2023 EC2 instance	Identity verification, bucket enumeration, ACL/policy and public-access queries
CloudGoat	2.5.0	Git clone; Python virtual environment; pip install	Deploying/working with an intentionally vulnerable AWS S3 training scenario
Steampipe	Installed during lab	Installed on EC2 and queried from shell	Querying S3 bucket public-access settings
jq	System package	Installed with dnf	Parsing JSON policy output and extracting principals

Tool / Service	Version	Installation Method	Purpose in This Task
Python	3.9.x	Amazon Linux package / virtual environment	Running the CloudGoat Python package
AWS Console	Web console	No local installation	IAM policy management and S3 configuration review

Key Configuration Commands Used

Key configuration and verification commands used:

```
python3 -m venv .venv
source .venv/bin/activate
pip install .
which cloudgoat
aws configure --profile cloudgoat
aws sts get-caller-identity --profile cloudgoat
cloudgoat list | grep -i s3
```

2.3 AI Tool Usage Declaration

AI Tool	Used For	What Was NOT Done with AI
ChatGPT	Explaining S3 policy concepts; troubleshooting Amazon Linux/CloudGoat setup; interpreting AccessDenied errors; helping structure the risk assessment and report language.	AI did not execute AWS commands, complete hosted labs, create lab screenshots, or independently generate technical evidence. Technical evidence came from the intern's authorised lab activity.

3.0 Methodology — Phase Execution

Phase A: Theory & Conceptual Foundation

Phase Objective	Understand how S3 exposure occurs through misconfigurations such as public ACLs, overly permissive bucket policies, or disabled Block Public Access settings.
Estimated Time	60 minutes
Evidence File(s)	phaseA_policy_analysis.png

Actions Taken

Reviewed the structure of an S3 bucket policy JSON document and identified the security significance of Effect, Principal, Action, and Resource. Analysed the meaning of a wildcard Principal ("*") and how it can make a statement applicable to anonymous/public principals when no effective blocking control prevents access. Reviewed the relationship between bucket policies, ACLs, IAM permissions, and S3 Block Public Access, and considered how attackers enumerate exposed bucket names and test accessible objects.

Key Observations

S3 exposure is usually caused by a combination of resource policy, ACL, IAM, and Block Public Access configuration rather than by a single setting. A wildcard Principal is a strong indicator requiring review because, when paired with Allow and sensitive S3 actions, it can create public exposure. Block Public Access provides an important preventive layer against accidental public configuration.

Evidence Produced

- phaseA_policy_analysis.png — screenshot documenting the Phase A S3 policy analysis and interpretation.

Phase B: Cloud Environment Setup & Target Preparation

Phase Objective	Deploy a misconfigured S3 environment using CloudGoat.
Estimated Time	90 minutes
Evidence File(s)	phaseB_s3_list.png

Actions Taken

Prepared a new Amazon Linux 2023 EC2 lab instance. Installed Git, Python/pip and supporting utilities with dnf, cloned CloudGoat, created and activated a Python virtual environment, and installed CloudGoat 2.5.0. Verified the CloudGoat executable, configured the AWS CLI profile named cloudgoat for the [REDACTED] IAM user, and verified identity with AWS STS. Listed S3-related CloudGoat scenarios and selected cloud_breach_s3 for the S3-focused exercise. Finally, enumerated the account buckets with `aws s3 ls --profile cloudgoat` to establish the environment baseline.

Key Observations

Amazon Linux 2023 uses dnf rather than apt, which required changing the initial installation commands. CloudGoat was installed as a Python package inside the virtual environment, so the `cloudgoat` command was used instead of `./cloudgoat.py`. CloudGoat also required an AWS profile/default configuration before scenario creation. The baseline S3 listing confirmed that the [REDACTED] credentials could enumerate bucket names.

Evidence Produced

- phaseB_s3_list.png — AWS CLI output showing the S3 buckets visible through the cloudgoat profile.

Phase C: Core Execution

Phase Objective	Enumerate all accessible buckets and analyze their ACLs/policies for exposure risks.
Estimated Time	150 minutes
Evidence File(s)	phaseC_public_buckets.png

Actions Taken

Enumerated accessible buckets using `aws s3 ls --profile cloudgoat` and `aws s3api list-buckets`. Tested bucket ACL retrieval with `aws s3api get-bucket-acl` and bucket policy retrieval with `aws s3api get-bucket-policy`. Initial ACL and

policy queries returned AccessDenied for several buckets because the [REDACTED] identity lacked the required read permissions. For the lab bucket secure-s3-lab-sk, policy retrieval showed a policy granting the [REDACTED] principal broad s3:* access to the bucket and its objects. S3 permissions were then consolidated for the lab user so the required public-access checks could be performed. Steampipe was used to query bucket public-access configuration with: select name, block_public_acls, block_public_policy from aws_s3_bucket;

Key Observations

Enumeration permissions and configuration-read permissions are separate: the user could list bucket names while still receiving AccessDenied for GetBucketAcl, GetBucketPolicy, and GetBucketPublicAccessBlock on resources not covered by its identity policy. This demonstrated least-privilege behavior in practice. Broad s3:* permissions simplified lab analysis but would be excessive for a production user. Public exposure assessment must consider both bucket policies/ACLs and Block Public Access settings.

Evidence Produced

- phaseC_public_buckets.png — Steampipe/AWS S3 public-access enumeration results used to assess bucket exposure.

Phase D: Independent Extension & Advanced Implementation

Phase Objective	Build a threat model showing how an attacker could exploit exposed S3 data.
Estimated Time	Approximately 120 minutes
Evidence File(s)	phaseD_risk_matrix

Actions Taken

Created an S3 risk assessment and threat model that mapped key assets, including data buckets, S3 objects, and logging buckets, to potential threats, security impacts, likelihood, and recommended mitigations. Risks were rated using **Low, Medium, and High** values, with controls such as S3 Block Public Access and least-privilege IAM permissions recommended.

Key Observations

Public S3 exposure and overly permissive IAM permissions can create significant confidentiality, integrity, and availability risks. Sensitive data exposure and unauthorized access were identified as **high-risk scenarios**, demonstrating the importance of preventing public access and limiting S3 permissions.

Evidence Produced

phaseD_risk_matrix.pdf : Matrix representation for risk.

Risk Matrix (Phase D Output)

Cloud Asset / Resource	Threat / Risk	MITRE Ref.	Likelihood (1–5)	Impact (1–5)	Score	Rating
S3 bucket policy / public access	Public exposure could permit unauthorised object access	T1530	4	5	20	Critical
IAM user [REDACTED]	Broad s3:* permission increases blast radius if credentials are misused	T1078	3	4	12	High
S3 configuration visibility	Insufficient permission/monitoring can create blind spots in exposure assessment	T1087	3	3	9	Medium
Public S3 metadata/account context	Enumeration may reveal useful reconnaissance information	T1087	3	3	9	Medium

Score = Likelihood × Impact | 1–4: Low | 5–9: Medium | 10–14: High | 15–25: Critical

Phase F: Documentation & Reporting

Phase Objective	Compile findings into a professional report demonstrating S3 exposure detection, policy analysis, enumeration, risk assessment, and remediation.
Estimated Time	60-90 minutes
Evidence File(s)	CSE_AWS_L_011_04_SK_20260829.pdf

Actions Taken

Compiled the technical work from Phases A-D into the EncryptEdge Labs task-report format. Reviewed the screenshots and command results, summarized the S3 exposure findings, mapped the risks to MITRE ATT&CK and defensive controls, documented remediation steps, and prepared the final report for PDF export. Sensitive identifiers and credentials are to be redacted from evidence before submission.

Key Observations

The combined evidence shows that S3 security depends on layered controls: IAM permissions determine what an authenticated principal can inspect, bucket policies and ACLs can introduce resource-level exposure, and Block Public Access can prevent or limit accidental public access. The troubleshooting performed during the lab also demonstrated that AccessDenied messages are useful evidence of effective authorization boundaries.

Evidence Produced

- CSE_AWS_L_011_04_SK_20260829.pdf — final task report containing methodology, findings, risk assessment, recommendations, reflection, and deliverables.

4.0 Findings

4.1 Findings Summary

#	Finding Title	Type / Category	Cloud Service / Resource	MITRE Ref.	Phase	Severity	Evidence Ref.
1	Potential S3 Public Exposure	Misconfiguration / Access Control Weakness	Amazon S3 bucket policies, ACLs, Block Public Access	T1530	Phase C	High	phaseC_public_buckets.png
2	Overly Broad S3 Permissions for Lab User	IAM Policy Gap / Access Control Weakness	IAM user [REDACTED] / S3 resources	T1078	Phase C	High	Bucket policy / IAM permission test screenshots
3	Enumeration Reveals Cloud Resource and Account Context	Information Exposure / Reconnaissance	Amazon S3 / hosted lab targets	T1087	Phase D	Medium	Cybr and PwnedLabs completion evidence
4	Authorization Gaps Produced AccessDenied During Assessment	Access Control / Visibility Gap	S3 ACL, policy and public-access APIs	N/A	Phase C	Low	Terminal AccessDenied screenshots

4.2 Detailed Finding Analysis

Finding 1: Potential S3 Public Exposure

Finding Title	Potential S3 Public Exposure
Type / Category	Misconfiguration / Access Control Weakness
Cloud Service / Resource	Amazon S3 bucket policies, ACLs and Block Public Access
Description	The assessment focused on whether bucket policies/ACLs and public-access-block settings could allow unintended public access. A wildcard Principal in an Allow statement would be a key exposure indicator. Steampipe and AWS CLI checks were used to inspect these controls.
How Identified	AWS CLI bucket-policy/ACL queries and Steampipe public-access configuration query.
MITRE ATT&CK Reference	T1530 - Data from Cloud Storage Object

Business / Security Impact	If public read access exists, unauthorised users could retrieve stored objects, affecting confidentiality and potentially exposing sensitive data.
Severity / Impact	High
Evidence Reference	Figure / evidence: phaseC_public_buckets.png

Insert the most relevant CLI output, JSON policy extract, or console screenshot for this finding below. Label it Figure X with a one-sentence caption.

Finding 2: Overly Broad S3 Permissions for Lab User

Finding Title	Overly Broad S3 Permissions for Lab User
Type / Category	IAM Policy Gap / Access Control Weakness
Cloud Service / Resource	IAM user [REDACTED] and S3 resources
Description	During troubleshooting, broad S3 permissions (s3:*) were used for the authorised lab user to permit required S3 configuration queries. This is acceptable for a temporary lab but is overly permissive for production use.
How Identified	IAM policy configuration and successful S3 policy/public-access queries after permission adjustment.
MITRE ATT&CK Reference	T1078 - Valid Accounts
Business / Security Impact	If [REDACTED] credentials were compromised while broad permissions remained, an attacker could read or alter multiple S3 resources, increasing confidentiality, integrity and availability impact.
Severity / Impact	High
Evidence Reference	Supporting IAM/S3 permission screenshots from Phase C

Insert the most relevant CLI output, JSON policy extract, or console screenshot for this finding below. Label it Figure X with a one-sentence caption.

Finding 3: Enumeration Reveals Cloud Resource and Account Context

Finding Title	Enumeration Reveals Cloud Resource and Account Context
Type / Category	Information Exposure / Reconnaissance
Cloud Service / Resource	Amazon S3 and authorised Cybr/PwnedLabs lab targets

Description	Bucket enumeration revealed that discovery can expose bucket names and access behavior even when detailed configuration actions are denied. Hosted labs further demonstrated how public S3 information can support account-context discovery.
How Identified	Cybr Introduction to S3 Enumeration and PwnedLabs Identify the AWS Account ID from a Public S3 Bucket.
MITRE ATT&CK Reference	T1087 - Account Discovery
Business / Security Impact	Reconnaissance information can help an attacker map cloud resources, identify promising targets, and prepare follow-on attempts against exposed data or identities.
Severity / Impact	Medium
Evidence Reference	Cybr and PwnedLabs lab completion screenshots

Insert the most relevant CLI output, JSON policy extract, or console screenshot for this finding below. Label it Figure X with a one-sentence caption.

5.0 Risk Assessment & Cloud Threat Model

5.1 Cloud Security Control Gaps & Defensive Coverage

Finding / Gap	Observable / Trigger in Cloud Telemetry	Recommended Control	AWS Service / Tool for Implementation
Potential S3 public exposure	Public/anonymous GetObject attempts; public-policy/ACL configuration changes	Enable account- and bucket-level Block Public Access; remove wildcard public grants; review ACLs	Amazon S3, AWS Config, IAM Access Analyzer, CloudTrail
Broad s3:* permission	Unusually broad S3 API activity by [REDACTED]; policy attachment/change events	Replace s3:* with least-privilege actions/resources after lab completion	AWS IAM, CloudTrail, IAM Access Analyzer
Enumeration/account context disclosure	Repeated ListBuckets/ListObjects and denied S3 configuration calls from unexpected identities/IPs	Limit enumeration permissions and alert on unusual S3 discovery patterns	CloudTrail, CloudWatch, Security Hub

5.2 Risk Exposure Over Time

Risk 1 — Public S3 Exposure: If a public policy or ACL remains uncorrected, unauthorised parties may continue discovering and retrieving objects. The exposure window grows over time and increases the likelihood of data exfiltration and compliance impact.

Risk 2 — Broad S3 Permissions: Leaving s3:* on a user after the lab creates unnecessary blast radius. Compromise of that valid account could permit broad S3 actions across resources accessible to the identity.

Risk 3 — Enumeration and Information Disclosure: Repeated resource discovery can provide attackers with bucket names, access behavior, and account context that supports follow-on targeting. Monitoring and least privilege reduce this reconnaissance value.

6.0 Recommendations

6.1 Technical Remediation

#	Finding	Specific Remediation Action (with Policy / CLI where applicable)	Priority	Standard / Reference
1	Potential S3 Public Exposure	Enable S3 Block Public Access at account/bucket level; remove any public ACL grants and wildcard Allow statements that are not explicitly required. Verify with <code>aws s3api get-public-access-block --bucket <bucket></code> .	Immediate	CIS AWS S3 public-access guidance; NIST SP 800-53 AC-3; ISO/IEC 27001 access control
2	Overly Broad S3 Permissions for Lab User	Remove temporary <code>s3:*</code> access after the lab and replace it with only the required List/Get configuration and object actions scoped to specific bucket ARNs.	Immediate	NIST SP 800-53 AC-6 Least Privilege; ISO/IEC 27001 access control
3	Enumeration Reveals Cloud Resource and Account Context	Restrict ListAllMyBuckets/ListBucket where not operationally required and review CloudTrail for unusual enumeration patterns.	Short-term	NIST SP 800-53 AU-6 / AC-6; MITRE T1087
4	S3 Configuration Monitoring	Enable AWS Config/IAM Access Analyzer/Security Hub checks for public S3 access and alert on risky policy/ACL changes.	Short-term	CIS AWS Foundations / AWS S3 security best practices

6.2 Strategic & Operational Improvements

- **Enforce S3 Block Public Access at the AWS account level** to prevent buckets from being unintentionally exposed through public ACLs or permissive bucket policies.
- **Implement least-privilege IAM permissions for S3 users and roles** by restricting broad permissions such as `s3:*` and granting only the actions and resources required for legitimate tasks.
- **Use AWS Config and IAM Access Analyzer for continuous S3 monitoring** to identify publicly accessible buckets, insecure bucket policies, ACL misconfigurations, and deviations from approved security configurations.
- **Strengthen S3 logging and monitoring** by enabling CloudTrail data events and protecting centralized audit logs so suspicious S3 activities, such as policy changes or unauthorized object access, can be investigated.
- **Perform periodic S3 security assessments** using AWS CLI and approved security tools to review bucket policies, ACLs, Block Public Access settings, encryption, and excessive permissions before they result in data exposure.

7.0 Reflection

1. How could an attacker leverage an exposed S3 bucket to escalate privileges within an AWS account?

An attacker could search an exposed S3 bucket for sensitive information such as AWS access keys, configuration files, scripts, backups, or application credentials. If valid credentials with greater permissions are discovered, the attacker could use them to access additional AWS resources and potentially escalate privileges. Overly permissive S3 write permissions could also allow an attacker to modify files used by applications or automated processes.

2. What business risks arise from storing sensitive data in publicly accessible cloud storage?

Publicly accessible S3 storage can lead to **data breaches, loss of confidentiality, regulatory violations, financial losses, and reputational damage**. Exposed customer or organizational information could be downloaded by unauthorized users, while writable buckets could also create data-integrity and availability risks.

3. How does shared responsibility between AWS and customers affect mitigation strategies?

AWS is responsible for **security of the cloud**, including the underlying infrastructure and physical facilities. The customer is responsible for **security in the cloud**, including S3 bucket policies, ACLs, IAM permissions, encryption choices, and public-access configuration. Therefore, organizations must correctly configure S3 security controls, apply least privilege, enable Block Public Access, and continuously monitor their AWS resources.

4. If you discovered such exposure in production, what would your incident response steps be?

I would first **contain the exposure** by blocking public access and removing unauthorized ACL or bucket-policy permissions. I would then preserve relevant logs and investigate CloudTrail and S3 activity to determine what data may have been accessed or modified. Any potentially compromised credentials would be disabled or rotated. After determining the scope and impact, I would remediate the underlying misconfiguration, restore affected data if necessary, notify appropriate security/compliance stakeholders, document the incident, and implement additional monitoring and preventive controls to reduce the chance of recurrence.

8.0 Deliverables Submission Checklist

✓	Phase	Deliverable Description	Your Filename	Status
✓	Phase A	Theory evidence - S3 policy structure and public-access analysis	phaseA_policy_analysis.png	Complete
✓	Phase B	CloudGoat/AWS environment setup and S3 bucket enumeration	phaseB_s3_list.png	Complete
✓	Phase C	Core S3 enumeration, policy/ACL and public-access analysis	phaseC_public_buckets.png	Complete

✓	Phase	Deliverable Description	Your Filename	Status
✓	Phase D	Build a threat model showing how an attacker could exploit exposed S3 data.	phaseD_risk_matrix	Complete
✓	Phase E	Hosted lab completion - PwnedLabs Identify the AWS Account ID from a Public S3 Bucket		Pending
✓	Phase E / Report	Final internship task report	CSE_AWS_L_011_04_SK_20260829.pdf	Complete

9.0 Additional Observations

Several troubleshooting points were important to the final result: Amazon Linux 2023 required dnf rather than apt; CloudGoat installed as a Python package and was invoked with the cloudgoat command; CloudGoat required an AWS profile before scenario creation; and S3 API/Steampipe checks initially returned AccessDenied because [REDACTED] lacked configuration-read permissions. These errors demonstrated the difference between bucket enumeration and permission to inspect ACL, policy, or Block Public Access settings. Temporary broad S3 permissions used for the lab should be removed after evidence collection.

- Unexpected AWS service behaviour encountered and how you resolved it
- CLI or tool errors and the troubleshooting steps taken — include error messages and resolution commands
- Cloud resource cost implications observed (e.g. unexpected charges in AWS Free Tier) and how you mitigated them
- Connections to other tasks or internship programme tracks you noticed during this task
- Optional premium labs or advanced extensions completed beyond the required Phase D tasks
- Feedback on the task brief, lab environment, or learning resources — the programme team reads this section

Declaration & Sign-Off

I confirm that:

- All work in this report is my own except where AI usage is explicitly declared in Section 2.3.
- All technical activities were performed exclusively within the authorised cloud environments listed in Section 1.3.
- No external AWS accounts, production environments, third-party cloud resources, or real customer data were accessed during this task.
- All AWS credentials, access keys, session tokens, and account identifiers used during this task have been redacted in evidence files before submission.
- I have reviewed this report for accuracy, completeness, and professional presentation before submission.

Intern Full Name	Signature	Date of Submission
[REDACTED]	[REDACTED]	29 August 2026

EncryptEdge Labs Limited — Junior Cloud Security Engineer Internship Programme

Company No. 15830711 | United Kingdom | encryptedgelabs.com

Copyright © 2026 EncryptEdge Labs. All rights reserved. | Report Template v3.0